

COMP5212 Machine Learning 2018 Fall programming project proposal

Building a self-driving agent with Reinforcement Learning

Hok Chun Ng, 20272532, hcngac@connect.ust.hk
Shengyuan Zhang, 20565161, szhangcg@connect.ust.hk
Ge Chen, 20360858, gchenaj@connect.ust.hk

October 15, 2018

Abstract

In this project proposal, we intend to build a self-driving agent that only relies on vision using reinforcement learning algorithms. Due to the limitation of real-world hardware and data, we intend to use a game simulator provided by OpenAI to generate training data. We introduce the elements of Q-learning, a reinforcement learning technique which does not require a model of the environment.

1 Description of the application and its practical significance

Self-driving car has been a hot topic in recent years. Several industry giants, like Google and Tesla Motor, have devoted significant efforts to developing self-driving cars. Letting computers to drive can not only release human beings from fatigue, but also reduce the frequency of traffic accidents. However, designing a robust self-driving system is non-trivial as the real-world traffic conditions are diversified. Due to the limitation of hardware and computing resources, we intend to find a solution to self-driving in a simulated game environment.

2 Formulation of the machine learning problems involved in the application

For human drivers, the driving behavior can be formulated in a loop from an environment sensing input to action output. Our eyes and ears are the sensors to interpret the current environment, including the road view from wind screen (e.g. weather, traffic lights, walking people), view from side mirror (e.g. neighboring cars, following cars), car horns, and so on. Then our brain takes all




Figure 1: Reinforcement Learning Illustration

these input signals to decide what action to take. The actions include turning the steering wheel, lighting turn signals, accelerating, braking, and so on. Once the actions are executed, the environment input changes, and again drivers will decide and take a new round of actions accordingly. This loop continues until the car arrives at the destination. Therefore, the machine learning problem in self-driving can be: given a set of environment input, find the best actions to execute until arriving at destination.

3 Data set

In this project, we intend to use the OpenAI Gym library to simulate car driving. The OpenAI Gym is a toolkit for developing and comparing reinforcement algorithms. It is a collection of test problems, by providing the environments that anyone can use to test her own reinforcement learning algorithms. The environments have a shared interface, allowing any user to write general algorithms.

We will focus on the environment “Enduro-v0”, which is the Atari 2600 game Enduro simulating a driving experience. In this environment, the observation is an RGB image of the screen (simulate the view from wind screen), which is an array of shape (210, 160, 3). Each action is repeatedly performed for a duration of k frames, where k is uniformly sampled from $\{2, 3, 4\}$. To reduce computation overhead, we may first converse the RGB images to grayscale ones.

4 Machine learning methods

The environment-action loop in the driving experience resembles the basics of reinforcement learning, an area of machine learning concerned with taking actions in an environment in order to maximize some notion of cumulative reward. Unlike supervised learning, reinforcement learning does not require correct input/output pairs. Instead, reinforcement learning often requires a reward function in terms of different actions under the environment. For driving, the final reward can be no traffic accident and arrive at the destination safely.

Figure 1 shows the basic model of reinforcement learning. At each time t , an agent receives the environment state S_t and current reward R_t . It then chooses an action a_t from the set of available actions, which is sent to the environment. The environment moves to a new state S_{t+1} and new reward R_{t+1} and feedback

to the agent. The goal of a reinforcement learning agent is to gain as much as reward as possible.

In this proposed project, we will use the method of deep Q-learning. Q-learning defines a reward expectation function, Q , which provides the expected reward of an action a taken under the current state s . This Q function can be recursively defined as $Q(s, a) = R(s, a) + \alpha \max_{a'} Q(S'(s, a), a')$, where R is the immediate reward given action a taken under state s , α is the discount rate of future reward and S' is the state transition. Given numerous possibilities of imagery state s , we cannot compute Q directly. Alternatively, we can use a deep convolutional network to estimate Q .

TODO reference insertion
TODO CNN description

5 Design of experiments and performance evaluation

TODO Experiment setup
TODO Performance evaluation
TODO Comparison with other model

6 Project planning

TODO project items
TODO project timeline
TODO work distribution

TODO References